

Multicollinearity in Linear Regression

```
In [21]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm
```

```
In [7]: df_adv = pd.read_csv(r"C:\Users\JANHAVI\Downloads\Advertising (1).csv", index_col=0)
X = df_adv[['TV', 'radio', 'newspaper']]
y = df_adv['sales']
df_adv.head()
```

```
Out[7]:
```

	TV	radio	newspaper	sales
1	230.1	37.8	69.2	22.1
2	44.5	39.3	45.1	10.4
3	17.2	45.9	69.3	9.3
4	151.5	41.3	58.5	18.5
5	180.8	10.8	58.4	12.9

```
In [8]: X = sm.add_constant(X)
```

```
In [9]: X
```

```
Out[9]:
```

	const	TV	radio	newspaper
1	1.0	230.1	37.8	69.2
2	1.0	44.5	39.3	45.1
3	1.0	17.2	45.9	69.3
4	1.0	151.5	41.3	58.5
5	1.0	180.8	10.8	58.4
...
196	1.0	38.2	3.7	13.8
197	1.0	94.2	4.9	8.1
198	1.0	177.0	9.3	6.4
199	1.0	283.6	42.0	66.2
200	1.0	232.1	8.6	8.7

200 rows × 4 columns

```
In [10]: ## fit a OLS model with intercept on TV and Radio

model = sm.OLS(y, X).fit() #OLS(endog = output feature, exog = input feature)
```

```
In [11]: model.summary()
```

Out[11]:

OLS Regression Results

Dep. Variable:	sales	R-squared:	0.897
Model:	OLS	Adj. R-squared:	0.896
Method:	Least Squares	F-statistic:	570.3
Date:	Mon, 01 Sep 2025	Prob (F-statistic):	1.58e-96
Time:	16:15:02	Log-Likelihood:	-386.18
No. Observations:	200	AIC:	780.4
Df Residuals:	196	BIC:	793.6
Df Model:	3		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	2.9389	0.312	9.422	0.000	2.324	3.554
TV	0.0458	0.001	32.809	0.000	0.043	0.049
radio	0.1885	0.009	21.893	0.000	0.172	0.206
newspaper	-0.0010	0.006	-0.177	0.860	-0.013	0.011

Omnibus:	60.414	Durbin-Watson:	2.084
Prob(Omnibus):	0.000	Jarque-Bera (JB):	151.241
Skew:	-1.327	Prob(JB):	1.44e-33
Kurtosis:	6.332	Cond. No.	454.

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [12]: import matplotlib.pyplot as plt
X.iloc[:,1:].corr()
```

Out[12]:

	TV	radio	newspaper
TV	1.000000	0.054809	0.056648
radio	0.054809	1.000000	0.354104
newspaper	0.056648	0.354104	1.000000

```
In [16]: df_salary = pd.read_csv(r"C:\Users\JANHAVI\Downloads\Salary_Data (2).csv")
df_salary.head()
```

Out[16]:

	YearsExperience	Age	Salary
0	1.1	21.0	39343
1	1.3	21.5	46205
2	1.5	21.7	37731
3	2.0	22.0	43525
4	2.2	22.2	39891

```
In [17]: X = df_salary[['YearsExperience', 'Age']]
y = df_salary['Salary']
```

```
In [18]: ## fit a OLS model with intercept on TV and Radio
X = sm.add_constant(X)
model = sm.OLS(y, X).fit()
```

```
In [19]: model.summary()
```

OLS Regression Results

Dep. Variable:	Salary	R-squared:	0.960
Model:	OLS	Adj. R-squared:	0.957
Method:	Least Squares	F-statistic:	323.9
Date:	Mon, 01 Sep 2025	Prob (F-statistic):	1.35e-19
Time:	16:17:22	Log-Likelihood:	-300.35
No. Observations:	30	AIC:	606.7
Df Residuals:	27	BIC:	610.9
Df Model:	2		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	-6661.9872	2.28e+04	-0.292	0.773	-5.35e+04	4.02e+04
YearsExperience	6153.3533	2337.092	2.633	0.014	1358.037	1.09e+04
Age	1836.0136	1285.034	1.429	0.165	-800.659	4472.686

Omnibus:	2.695	Durbin-Watson:	1.711
Prob(Omnibus):	0.260	Jarque-Bera (JB):	1.975
Skew:	0.456	Prob(JB):	0.372
Kurtosis:	2.135	Cond. No.	626.

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [20]: X.iloc[:,1:].corr()
```

Out[20]:

	YearsExperience	Age
YearsExperience	1.000000	0.987258
Age	0.987258	1.000000