

# Ridge & Lasso Regression

```
In [20]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
from sklearn.linear_model import Lasso, Ridge
```

```
In [41]: import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: from sklearn.datasets import fetch_california_housing
```

```
In [3]: df = fetch_california_housing()
df
```

```
Out[3]: {'data': array([[ 8.3252, 41., 6.98412698, ..., 2.55555556,
        37.88, -122.23],
       [ 8.3014, 21., 6.23813708, ..., 2.10984183,
        37.86, -122.22],
       [ 7.2574, 52., 8.28813559, ..., 2.80225989,
        37.85, -122.24],
       ...,
       [ 1.7, 17., 5.20554273, ..., 2.3256351,
        39.43, -121.22],
       [ 1.8672, 18., 5.32951289, ..., 2.12320917,
        39.43, -121.32],
       [ 2.3886, 16., 5.25471698, ..., 2.61698113,
        39.37, -121.24]]),
  'target': array([4.526, 3.585, 3.521, ..., 0.923, 0.847, 0.894]),
  'frame': None,
  'target_names': ['MedHouseVal'],
  'feature_names': ['MedInc',
    'HouseAge',
    'AveRooms',
    'AveBedrms',
    'Population',
    'AveOccup',
    'Latitude',
    'Longitude'],
  'DESCR': '.. _california_housing_dataset:\n\nCalifornia Housing dataset\n-----\n\n\n**Data Set Characteristics:**\n\n\nNumber of Instances: 20640\n\n\nNumber of Attributes: 8 numeric, predictive attributes and the target\n\n\nAttribute Information:\n    - MedInc          median income in block group\n    - HouseAge       median house age in block group\n    - AveRooms        average number of rooms per household\n    - AveBedrms       average number of bedrooms per household\n    - Population      block group population\n    - AveOccup        average number of household members\n    - Latitude        block group latitude\n    - Longitude       block group longitude\n\n\nMissing Attribute Values: None\n\n\nThis dataset was obtained from the StatLib repository.\nhttps://www.dcc.fc.up.pt/~ltorgo/Regression/cal_housing.html\n\n\nThe target variable is the median house value for California districts,\nexpressed in hundreds of thousands of dollars ($100,000).\n\n\nThis dataset was derived from the 1990 U.S. census, using one row per census\nblock group. A block group is the smallest geographical unit for which the U.S.\nCensus Bureau publishes sample data (a block group typically has a population\nof 600 to 3,000 people).\n\n\nA household is a group of people residing within a home. Since the average\nnumber of rooms and bedrooms in this dataset are provided per household, these\ncolumns may take surprisingly large values for block groups with few households\nand many empty houses, such as vacation resorts.\n\n\nIt can be downloaded/loaded using the\nfunc:`sklearn.datasets.fetch_california_housing` function.\n\n\n.. rubric:: References\n\n- Pace, R. Kelley and Ronald Barry, Sparse Spatial Autoregressions,\nStatistics and Probability Letters, 33:291-297, 1997.\n'}
```

```
In [4]: housing = fetch_california_housing(as_frame=True)
```

```
In [5]: df = housing.frame
```

```
In [6]: x = housing.data      #independent Variable
        y = housing.target    #dependent variable
```

```
In [7]: df.head()
```

Out[7]:

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	Longitude	MedInc
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556	37.88	-122.23	
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842	37.86	-122.22	
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260	37.85	-122.24	
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945	37.85	-122.25	
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467	37.85	-122.25	

In [8]: `df.shape`

Out[8]: (20640, 9)

In [9]: `df.head()`

Out[9]:

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	Longitude	MedInc
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556	37.88	-122.23	
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842	37.86	-122.22	
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260	37.85	-122.24	
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945	37.85	-122.25	
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467	37.85	-122.25	

In [10]: `X= df.iloc[:, :-1]`  
`Y=df.iloc[:, -1]`

## Linear Regression

In [11]: `from sklearn.model_selection import cross_val_score`  
`from sklearn.linear_model import LinearRegression`

```
lin_regressor=LinearRegression()
mse=cross_val_score(lin_regressor,X,Y,scoring='neg_mean_squared_error',cv=5)
mean_mse=np.mean(mse)
print(mean_mse)
```

-0.558290171768654


## Ridge Regression

In [12]: `# FIND THE VALUES OF LAMBDA & LAMBDA USED TO FIND A CROSS VALIDATION`

In [13]: `from sklearn.linear_model import Ridge`  
`from sklearn.model_selection import GridSearchCV`

```
ridge=Ridge()
parameters={'alpha':[1e-15,1e-10,1e-8,1e-3,1e-2,1,5,10,20,30,35,40,45,50,55,100]}
ridge_regressor=GridSearchCV(ridge,parameters,scoring='neg_mean_squared_error',cv=5)
ridge_regressor.fit(X,Y)
```

Out[13]:



In [14]:

```
print(ridge_regressor.best_params_)
print(ridge_regressor.best_score_)# Mean Squared Error

{'alpha': 55}
-0.5579444917053032
```


## Lasso Regression

In [15]:

```
from sklearn.linear_model import Lasso
from sklearn.model_selection import GridSearchCV

lasso=Lasso()
parameters={'alpha':[1e-15,1e-10,1e-8,1e-3,1e-2,1,5,10,20,30,35,40,45,50,55,100]}
Lasso_regressor=GridSearchCV(lasso,parameters,scoring='neg_mean_squared_error',cv=5)
Lasso_regressor.fit(X,Y)
```

Out[15]:



In [16]:

```
print(Lasso_regressor.best_params_)
print(Lasso_regressor.best_score_)

{'alpha': 55}
-0.5579444917053032
```

In [26]:

```
from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(X,Y, test_size=0.3, random_state=42)
from sklearn.linear_model import Lasso, Ridge
from sklearn.metrics import mean_squared_error, r2_score
```

In [32]:

```
from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.linear_model import Lasso, Ridge
import pandas as pd

# Split dataset
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Step 1: Create model
lasso_regressor = Lasso(alpha=0.1)
ridge_regressor = Ridge(alpha=1.0)

# Step 2: Train model
lasso_regressor.fit(X_train, y_train)
```


```
ridge_regressor.fit(X_train, y_train)

# Step 3: Predict
prediction_lasso = lasso_regressor.predict(X_test)
prediction_ridge = ridge_regressor.predict(X_test)
```

```
In [33]: prediction_lasso=lasso_regressor.predict(X_test)
         prediction_ridge=ridge_regressor.predict(X_test)
```


```
In [42]: sns.distplot(Y_test-prediction_lasso)
```

```
Out[42]: <Axes: xlabel='MedHouseVal', ylabel='Density'>
```




```
In [39]: sns.histplot(Y_test - prediction_lasso, kde=True)
```

```
Out[39]: <Axes: xlabel='MedHouseVal', ylabel='Count'>
```



```
In [49]: sns.distplot(Y_test-prediction_ridge)
```

```
Out[49]: <Axes: xlabel='MedHouseVal', ylabel='Density'>
```



```
In [48]: sns.histplot(Y_test - prediction_ridge, kde=True)
```

```
Out[48]: <Axes: xlabel='MedHouseVal', ylabel='Count'>
```

